

GITHUB LINK :

https://github.com/Pujya6267/21BDS0372_EDA_Theory_DA

EXPLORATORY DATA ANALYSIS

THEORY DIGITAL ASSIGNMENT

PUJYA JAIN

21BDS0372

Question) Procedure to Follow:

1. Load the Dataset and explore Dimension, Summary, Data Handling, Data Cleaning, Univariate, Bivariate and Multivariate Analysis (Maximum of all the possibilities).
2. do statistical analysis on the data.
3. Estimation of missing data, global methods, class-based methods, multiple imputation methods etc.,
4. Use statistical techniques to identify outlier data for the given data set.
5. find covariance and correlation using functions.
6. and as many more operations as feels relevant.

1) Load the Data and Basic Exploration

```
In [30]: import pandas as pd

data = pd.read_csv('USSeatBelts.csv')
print("\nFirst Five Rows:\n", data.head())
```

First Five Rows:

	rownames	state	year	miles	fatalities	seatbelt	speed65	speed70	drinkage	\
0	1	AK	1983	3358	0.044669	NaN	no	no	yes	
1	2	AK	1984	3589	0.037336	NaN	no	no	yes	
2	3	AK	1985	3840	0.033073	NaN	no	no	yes	
3	4	AK	1986	4008	0.025200	NaN	no	no	yes	
4	5	AK	1987	3900	0.019487	NaN	no	no	yes	

	alcohol	income	age	enforce
0	no	17973	28.234966	no
1	no	18093	28.343542	no
2	no	18925	28.372816	no
3	no	18466	28.396652	no
4	no	18021	28.453251	no

```
In [31]: print("Data Shape:", data.shape)
```

Data Shape: (765, 13)

```
In [32]: print("Column Names:", data.columns)
```

```
Column Names: Index(['rownames', 'state', 'year', 'miles', 'fatalities', 'seatbelt',  
                    'speed65', 'speed70', 'drinkage', 'alcohol', 'income', 'age',  
                    'enforce'],  
                    dtype='object')
```

```
In [33]: print("\nData Types:\n", data.dtypes)
```

```
Data Types:  
rownames      int64  
state         object  
year          int64  
miles         int64  
fatalities    float64  
seatbelt      float64  
speed65       object  
speed70       object  
drinkage      object  
alcohol       object  
income        int64  
age           float64  
enforce       object  
dtype: object
```

```
In [34]: print("\nMissing Values:\n", data.isnull().sum())
```

```
Missing Values:  
rownames      0  
state         0  
year          0  
miles         0  
fatalities    0  
seatbelt     209  
speed65       0  
speed70       0  
drinkage      0  
alcohol       0  
income        0  
age           0  
enforce       0  
dtype: int64
```

2)Data Cleaning and Imputation We found that seatbelt had missing values, which we imputed with the mean:

```
In [35]: # Calculate missing values percentage  
missing_percentage = (data.isnull().sum() / len(data)) * 100  
print("Missing Values (Percentage):\n", missing_percentage)
```

```
Missing Values (Percentage):  
rownames      0.000000  
state         0.000000  
year          0.000000  
miles         0.000000  
fatalities    0.000000  
seatbelt     27.320261  
speed65       0.000000  
speed70       0.000000  
drinkage      0.000000  
alcohol       0.000000  
income        0.000000  
age           0.000000  
enforce       0.000000  
dtype: float64
```

```
In [61]: # Impute missing values in 'seatbelt' using mean as an example  
# Create a copy of the DataFrame to avoid chained assignment issues  
data = data.copy()
```

```
# Impute missing values in 'seatbelt' using the mean
data['seatbelt'] = data['seatbelt'].fillna(data['seatbelt'].mean())

# Check for remaining missing values
print(data.isnull().sum())
```

```
rownames      0
state         0
year          0
miles         0
fatalities    0
seatbelt      0
speed65       0
speed70       0
drinkage      0
alcohol       0
income        0
age           0
enforce       0
miles_log     0
dtype: int64
```

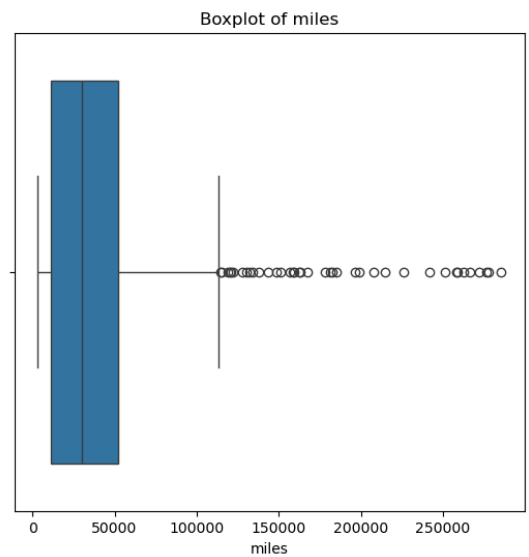
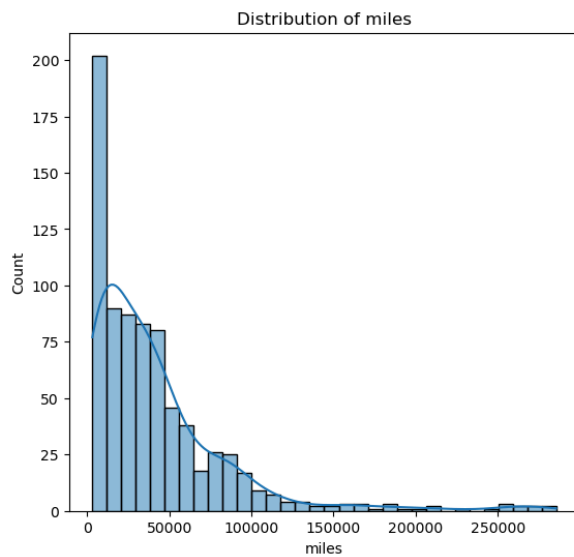
3. Univariate Analysis with Additional Statistical Details

```
In [37]: import matplotlib.pyplot as plt
import seaborn as sns

numerical_columns = ['miles', 'fatalities', 'seatbelt', 'income', 'age']

for column in numerical_columns:
    print(f"\n{column} Summary:")
    print(f"Mean: {data[column].mean()}")
    print(f"Median: {data[column].median()}")
    print(f"Skewness: {data[column].skew()}")
    print(f"Kurtosis: {data[column].kurt()}")
    print(f"Unique Values: {data[column].nunique()}")
    plt.figure(figsize=(14, 6))
    plt.subplot(1, 2, 1)
    sns.histplot(data[column], kde=True)
    plt.title(f'Distribution of {column}')
    plt.subplot(1, 2, 2)
    sns.boxplot(x=data[column])
    plt.title(f'Boxplot of {column}')
    plt.show()
```

```
miles Summary:
Mean: 41447.734640522875
Median: 30319.0
Skewness: 2.6111023792585244
Kurtosis: 9.063223045992348
Unique Values: 758
```



fatalities Summary:

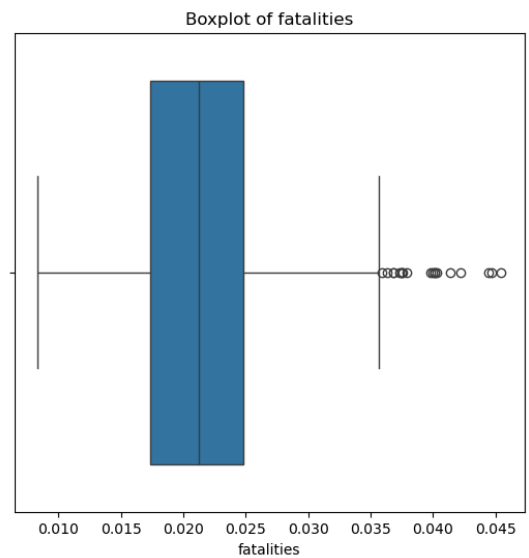
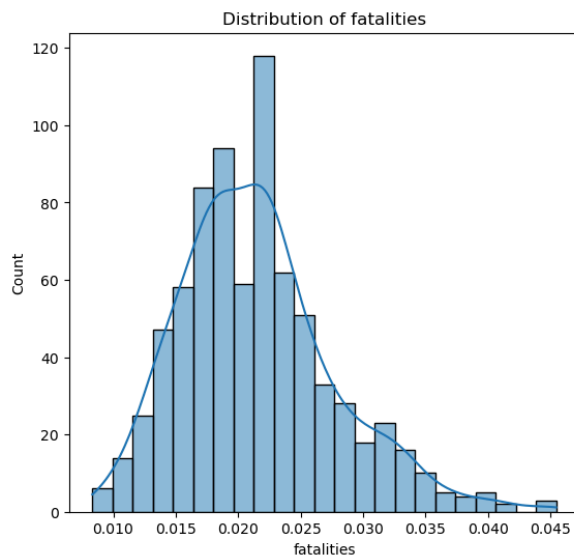
Mean: 0.02148951477572025

Median: 0.021198958158493

Skewness: 0.7600253235279069

Kurtosis: 0.8113551489655251

Unique Values: 765



seatbelt Summary:

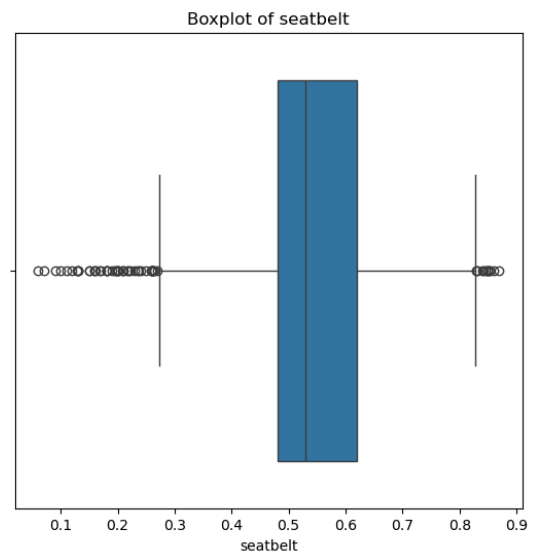
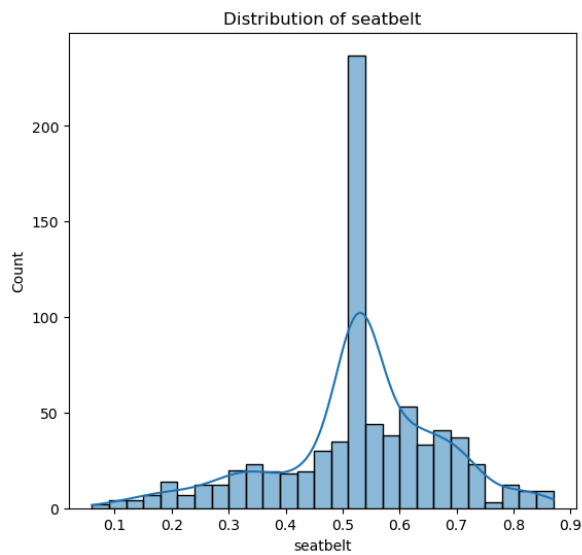
Mean: 0.528851797207922

Median: 0.528851797207922

Skewness: -0.5507606490101725

Kurtosis: 0.6626158077152864

Unique Values: 246



income Summary:

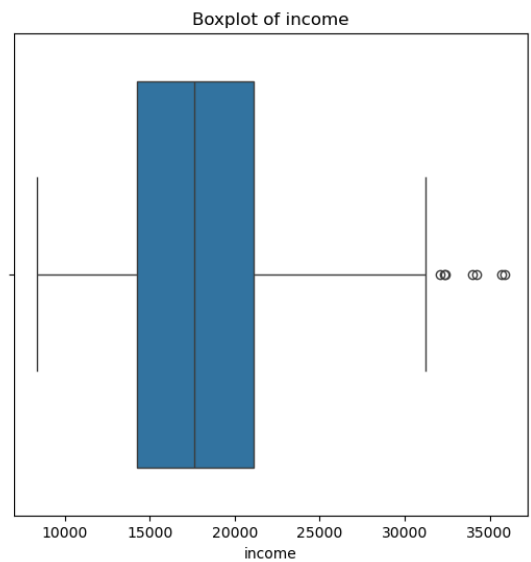
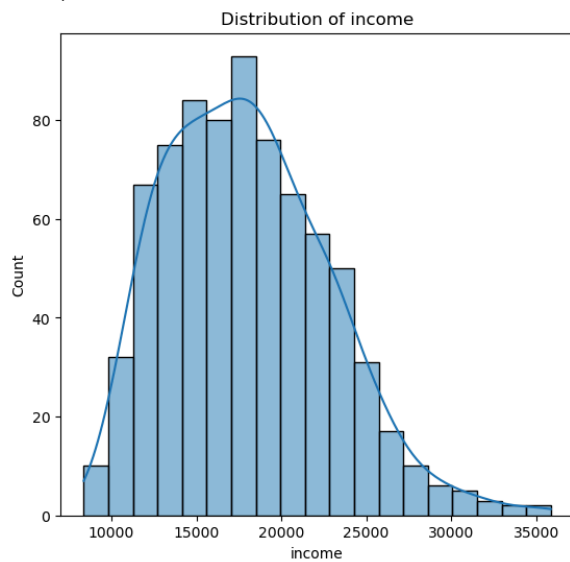
Mean: 17992.586928104574

Median: 17624.0

Skewness: 0.5777915492689545

Kurtosis: 0.2105946820100839

Unique Values: 750



age Summary:

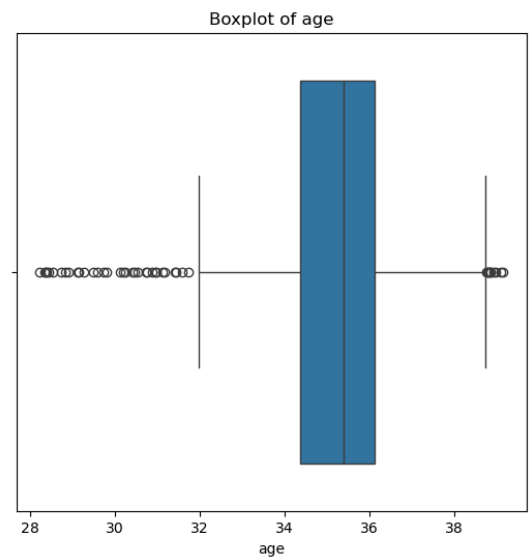
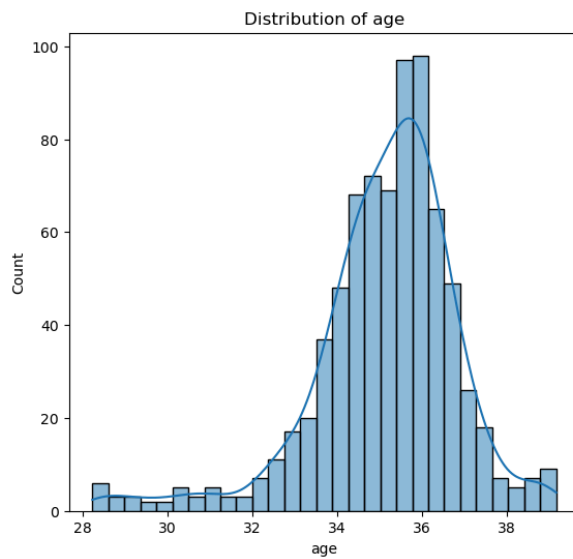
Mean: 35.13719353270687

Median: 35.39176559448242

Skewness: -1.1495976096228606

Kurtosis: 2.9703837051798776

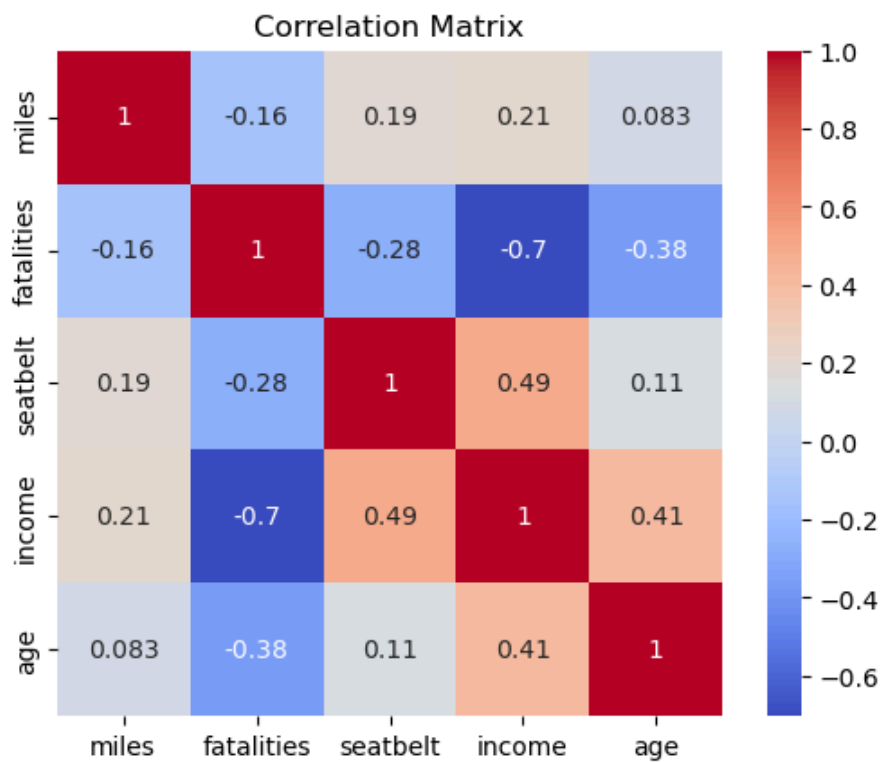
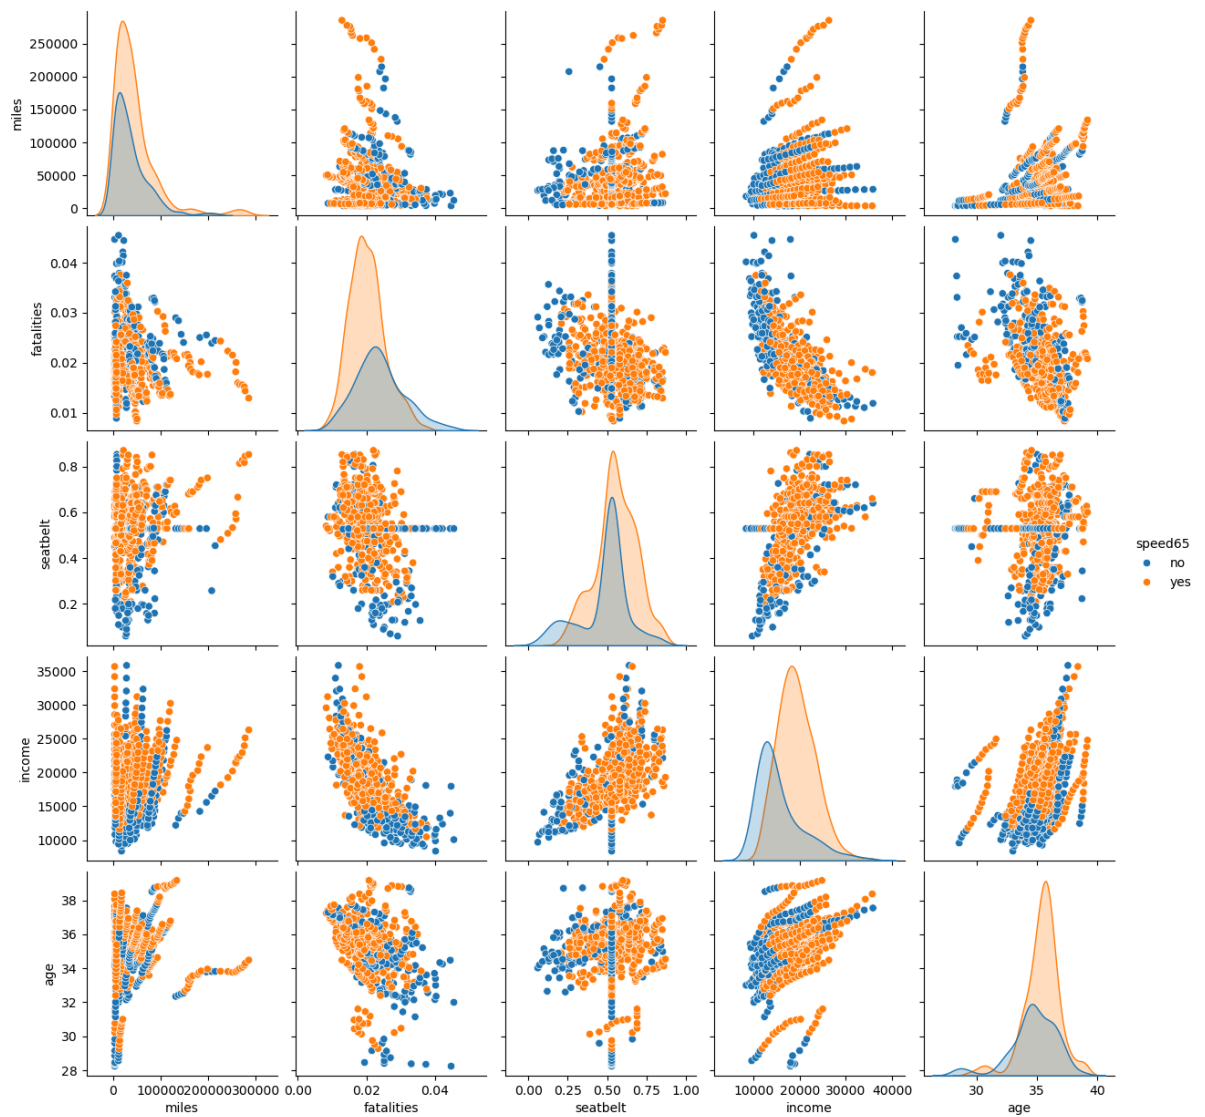
Unique Values: 764



4. Bivariate and Multivariate Analysis

```
In [38]: # Pairplot for bivariate analysis with 'speed65' as a hue
sns.pairplot(data, vars=numerical_columns, hue='speed65')
plt.show()

# Correlation matrix heatmap
correlation_matrix = data[numerical_columns].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', square=True)
plt.title("Correlation Matrix")
plt.show()
```



5. Outlier Detection Using Multiple Methods

In [39]: **from** scipy **import** stats

```
# IQR method for outlier detection
def detect_outliers_iqr(column):
    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)
    IQR = Q3 - Q1
    outliers = column[((column < (Q1 - 1.5 * IQR)) | (column > (Q3 + 1.5 * IQR)))]
    return outliers

for column in numerical_columns:
    print(f"\nOutliers in {column} using IQR method:")
    print(detect_outliers_iqr(data[column]))

# Z-score method for outlier detection
z_scores = stats.zscore(data[numerical_columns])
abs_z_scores = abs(z_scores)
outliers_z = (abs_z_scores > 3).any(axis=1) # Outliers where any Z-score > 3
print(f"\nNumber of Outliers (Z-score method):", outliers_z.sum())
```


Outliers in miles using IQR method:

60	182652
61	196537
62	207600
63	214913
64	226301
65	241575
66	251482
67	258926
68	257976
69	262548
70	266408
71	271943
72	276371
73	278043
74	285612
144	114311
145	120467
146	121989
147	127801
148	130004
149	134007
522	115091
523	118641
524	120778
645	131883
646	137737
647	143263
648	148348
649	151186
650	156458
651	159512
652	162232
653	158756
654	163329
655	167611
656	178348
657	181096
658	185386
659	198700

Name: miles, dtype: int64

Outliers in fatalities using IQR method:

0	0.044669
1	0.037336
46	0.042158
47	0.041381
48	0.044430
375	0.040164
376	0.036818
378	0.040102
379	0.037476
390	0.039827
480	0.045470
481	0.039977
482	0.040320
483	0.037886
484	0.037576
495	0.036816
603	0.037487
604	0.035932
735	0.036337

Name: fatalities, dtype: float64

Outliers in seatbelt using IQR method:

16	0.130
17	0.170

```
19      0.210
33      0.198
62      0.258
...
692     0.181
693     0.218
718     0.840
724     0.263
757     0.260
Name: seatbelt, Length: 64, dtype: float64
```

Outliers in income using IQR method:

```
102     32073
103     33979
104     35863
117     32398
118     34213
119     35704
479     32356
Name: income, dtype: int64
```

Outliers in age using IQR method:

```
0      28.234966
1      28.343542
2      28.372816
3      28.396652
4      28.453251
5      28.851419
6      29.148954
7      29.586285
8      29.827711
9      30.210697
10     30.464386
11     30.756571
12     31.178596
13     31.445354
14     31.601475
139    38.802761
140    38.823399
141    38.874424
143    38.766312
144    38.835045
145    38.864780
146    38.926968
147    38.998173
148    39.102768
149    39.169582
660    28.554277
661    28.735720
662    28.915110
663    29.123068
664    29.273216
665    29.507109
666    29.745262
667    30.124149
668    30.248541
669    30.431841
670    30.539591
671    30.742689
672    30.891895
673    30.955338
674    30.997065
750    31.140636
751    31.433319
752    31.746651
Name: age, dtype: float64
```

Number of Outliers (Z-score method): 49

6. Covariance and Correlation Calculations with Interpretation

```
In [40]: # Covariance Matrix
covariance_matrix = data[numerical_columns].cov()
print("Covariance Matrix:\n", covariance_matrix)

# Correlation Matrix
correlation_matrix = data[numerical_columns].corr()
print("Correlation Matrix:\n", correlation_matrix)

# Interpretation: Look for strong correlations
print("\nHighly Correlated Pairs (Threshold > 0.6):")
high_corr = correlation_matrix[(correlation_matrix > 0.6) & (correlation_matrix < 1)]
print(high_corr.dropna(how='all').dropna(axis=1, how='all'))
```

Covariance Matrix:

	miles	fatalities	seatbelt	income	age
miles	1.932657e+09	-43.779000	1188.702267	4.360538e+07	6230.376557
fatalities	-4.377900e+01	0.000038	-0.000250	-2.089069e+01	-0.003934
seatbelt	1.188702e+03	-0.000250	0.021040	3.403379e+02	0.028292
income	4.360538e+07	-20.890689	340.337944	2.315014e+07	3329.698388
age	6.230377e+03	-0.003934	0.028292	3.329698e+03	2.883650

Correlation Matrix:

	miles	fatalities	seatbelt	income	age
miles	1.000000	-0.161366	0.186411	0.206151	0.083458
fatalities	-0.161366	1.000000	-0.279736	-0.703558	-0.375413
seatbelt	0.186411	-0.279736	1.000000	0.487652	0.114860
income	0.206151	-0.703558	0.487652	1.000000	0.407527
age	0.083458	-0.375413	0.114860	0.407527	1.000000

Highly Correlated Pairs (Threshold > 0.6):

Empty DataFrame

Columns: []

Index: []

7. Additional Operations

Skewness and Transformation: Apply transformations if features are skewed (e.g., log transformation for positive skew).

Categorical Encoding: Convert categorical variables to numerical values for further analysis.

Feature Engineering: Create new features based on existing columns, such as ratios or interaction terms.

```
In [41]: # Skewed feature transformation (example: Log transformation for 'miles')
import numpy as np

if data['miles'].skew() > 0.5:
    data['miles_log'] = np.log1p(data['miles']) # log1p handles zero values better
    print("Applied log transformation on 'miles'")

# Categorical encoding
data_encoded = pd.get_dummies(data, columns=['speed65', 'speed70', 'drinkage', 'alcohol', ''])

print("Encoded Data Columns:\n", data_encoded.columns)
```

Applied log transformation on 'miles'

Encoded Data Columns:

```
Index(['rownames', 'state', 'year', 'miles', 'fatalities', 'seatbelt',  
      'income', 'age', 'miles_log', 'speed65_no', 'speed65_yes', 'speed70_no',  
      'speed70_yes', 'drinking_no', 'drinking_yes', 'alcohol_no',  
      'alcohol_yes', 'enforce_no', 'enforce_primary', 'enforce_secondary'],  
      dtype='object')
```

8. Advanced Imputation Techniques

In [44]: `from scipy.stats import ttest_ind`

```
# Example: T-test between 'speed65' categories for 'fatalities'  
fatalities_yes = data[data['speed65'] == 'yes']['fatalities']  
fatalities_no = data[data['speed65'] == 'no']['fatalities']  
  
t_stat, p_val = ttest_ind(fatalities_yes, fatalities_no)  
print(f"T-test between speed65 'yes' and 'no': T-Statistic = {t_stat}, P-Value = {p_val}")
```

T-test between speed65 'yes' and 'no': T-Statistic = -8.113936105023495, P-Value = 1.956756082051727e-15

9. Data Preprocessing for applying ML models

In [46]: `from sklearn.model_selection import train_test_split`
`from sklearn.preprocessing import LabelEncoder, StandardScaler`

```
# Drop rows with missing target column ('seatbelt') for supervised models  
data_cleaned = data.dropna(subset=['seatbelt'])  
  
# Separate numerical and categorical features  
numerical_features = ['miles', 'fatalities', 'income', 'age']  
categorical_features = ['state', 'speed65', 'speed70', 'drinking', 'alcohol', 'enforce']  
  
# Encode categorical features  
encoder = LabelEncoder()  
for col in categorical_features:  
    data_cleaned[col] = encoder.fit_transform(data_cleaned[col])  
  
# Scale numerical features  
scaler = StandardScaler()  
data_cleaned[numerical_features] = scaler.fit_transform(data_cleaned[numerical_features])  
  
# Define features (X) and target (y)  
X = data_cleaned[numerical_features + categorical_features]  
y = data_cleaned['seatbelt']  
  
# Split into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
# Check preprocessed data  
X_train.head(), y_train.head()
```

```
Out[46]: (      miles  fatalities  income      age  state  speed65  speed70  \
336  0.875015   -0.157899 -0.001162 -0.216020     22        1        0
688  0.679537   -1.490081  1.446955  0.423421     45        1        0
290  0.042935   -0.771656  0.712186  0.656757     19        0        0
687  0.645599   -1.394065  1.237526  0.313560     45        1        0
90   -0.473849   -0.041882 -0.493434  0.528608      6        0        0

      drinkage  alcohol  enforce
336          1         0         2
688          1         1         2
290          1         0         0
687          1         1         2
90           0         0         0 ,
336    0.467000
688    0.696000
290    0.528852
687    0.702000
90     0.528852
Name: seatbelt, dtype: float64)
```

LINEAR REGRESSION

```
In [52]: from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Linear Regression
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
y_pred_lr = lin_reg.predict(X_test)

{
    "Linear Regression": {"MSE": lr_mse, "R^2": lr_r2}
}
```

```
Out[52]: {'Linear Regression': {'MSE': 0.01606315012019666, 'R^2': 0.1408500034477176}}
```

RANDOM FOREST REGRESSOR

```
In [54]: # Random Forest Regressor
rf_reg = RandomForestRegressor(random_state=42)
rf_reg.fit(X_train, y_train)
y_pred_rf = rf_reg.predict(X_test)

# Evaluate models
lr_mse = mean_squared_error(y_test, y_pred_lr)
lr_r2 = r2_score(y_test, y_pred_lr)

rf_mse = mean_squared_error(y_test, y_pred_rf)
rf_r2 = r2_score(y_test, y_pred_rf)

{
    "Random Forest Regressor": {"MSE": rf_mse, "R^2": rf_r2}
}
```

```
Out[54]: {'Random Forest Regressor': {'MSE': 0.008952132765534183,
    'R^2': 0.5211882615120511}}
```

RANDOM FOREST CLASSIFIER

```
In [59]: from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import StandardScaler
```

```

log_reg = LogisticRegression(random_state=42, max_iter=500)
log_reg.fit(X_train, y_train_class)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

log_reg = LogisticRegression(random_state=42)
log_reg.fit(X_train_scaled, y_train_class)
y_pred_lr_class = log_reg.predict(X_test_scaled)

log_reg = LogisticRegression(random_state=42, solver='saga', max_iter=500)
log_reg.fit(X_train, y_train_class)

{
    "Logistic Regression Accuracy": lr_accuracy
}

```

C:\Users\harsh\anaconda3\Lib\site-packages\sklearn\linear_model_sag.py:349: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(

Out[59]: {'Logistic Regression Accuracy': 0.7908496732026143}

LOGISTIC REGRESSION

```

In [57]: # Random Forest Classifier
rf_clf = RandomForestClassifier(random_state=42)
rf_clf.fit(X_train, y_train_class)
y_pred_rf_class = rf_clf.predict(X_test)

# Evaluate models
lr_accuracy = accuracy_score(y_test_class, y_pred_lr_class)
rf_accuracy = accuracy_score(y_test_class, y_pred_rf_class)

{
    "Random Forest Classifier Accuracy": rf_accuracy,
    "Classification Report (Random Forest)": classification_report(y_test_class, y_pred_rf_
}

```

Out[57]: {'Random Forest Classifier Accuracy': 0.8954248366013072,
'Classification Report (Random Forest)': {'0': {'precision': 0.9333333333333333,
'recall': 0.8936170212765957,
'f1-score': 0.9130434782608695,
'support': 94.0},
'1': {'precision': 0.8412698412698413,
'recall': 0.8983050847457628,
'f1-score': 0.8688524590163934,
'support': 59.0},
'accuracy': 0.8954248366013072,
'macro avg': {'precision': 0.8873015873015873,
'recall': 0.8959610530111792,
'f1-score': 0.8909479686386315,
'support': 153.0},
'weighted avg': {'precision': 0.8978317252827056,
'recall': 0.8954248366013072,
'f1-score': 0.8960024969835878,
'support': 153.0}}}

NAIVE BAYES CLASSIFIER

```

In [60]: from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report

```

```

# Initialize the Gaussian Naive Bayes model
nb_model = GaussianNB()

# Fit the model on the training data
nb_model.fit(X_train_scaled, y_train_class)

# Predict on the test data
y_pred_nb = nb_model.predict(X_test_scaled)

# Evaluate the model
nb_accuracy = accuracy_score(y_test_class, y_pred_nb)
nb_report = classification_report(y_test_class, y_pred_nb)

print("Naive Bayes Accuracy:", nb_accuracy)
print("\nClassification Report for Naive Bayes:\n", nb_report)

```

Naive Bayes Accuracy: 0.6143790849673203

Classification Report for Naive Bayes:

	precision	recall	f1-score	support
0	1.00	0.37	0.54	94
1	0.50	1.00	0.67	59
accuracy			0.61	153
macro avg	0.75	0.69	0.60	153
weighted avg	0.81	0.61	0.59	153